


Review of Lecture 2

Is Learning feasible?

Yes, in a **probabilistic** sense.



$$\mathbb{P} \left[|E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon \right] \leq 2e^{-2\epsilon^2 N}$$

Since g has to be one of h_1, h_2, \dots, h_M , we conclude that

If:

$$|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon$$

Then:

$$|E_{\text{in}}(h_1) - E_{\text{out}}(h_1)| > \epsilon \quad \text{or}$$

$$|E_{\text{in}}(h_2) - E_{\text{out}}(h_2)| > \epsilon \quad \text{or}$$

...

$$|E_{\text{in}}(h_M) - E_{\text{out}}(h_M)| > \epsilon$$

The problem with having multiple hypotheses is that there is a possible scenario where the probability of a hypothesis generalising badly (0.5% for h_1 , 0.5% for h_2 , etc.) accumulates and we end up with a significant probability that one of the hypotheses is bad. If we are unlucky that our final chosen hypothesis is this hypothesis, E_{in} will not track E_{out} for this chosen hypothesis. We use this to say that if g generalises badly, at least one of the hypotheses we tested must also be bad, hence the if-then statement. Applying the union bound we can simplify this inequality - eventually ending up with the M factor.

This gives us an added **M** factor.

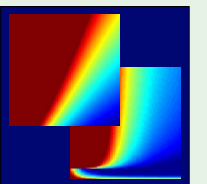
Learning From Data

Yaser S. Abu-Mostafa
California Institute of Technology

Lecture 3: **Linear Models I**



Sponsored by Caltech's Provost Office, E&AS Division, and IST • Tuesday, April 10, 2012



Outline

- Input representation
- Linear Classification
- Linear Regression
- Nonlinear Transformation

A real data set



Input representation

'raw' input $\mathbf{x} = (x_0, x_1, x_2, \dots, x_{256})$

linear model: $(w_0, w_1, w_2, \dots, w_{256})$

If there are too many parameters, the model is unlikely to generalize well, so need to extract useful information (ideally when we do feature extraction, the data we lose is largely irrelevant)

Features: Extract useful information, e.g.,

intensity and symmetry $\mathbf{x} = (x_0, x_1, x_2)$

linear model: (w_0, w_1, w_2)





Illustration of features

$$\mathbf{x} = (x_0, x_1, x_2)$$


x_1 : intensity

x_2 : symmetry




What PLA does

Evolution of E_{in} and E_{out}




Final perceptron boundary



Typical behaviour of PLA since data not linearly separable - also it will not converge, so forced to stop at 1000. Note that E_{out} tracks E_{in} well since there are only 3 parameters (very simple model) so will generalize well


The 'pocket' algorithm

PLA:




Pocket:

Simple modification where we keep track of the in-sample error E_{in} and keep track of it throughout if it is not an improvement, skip and carry on with the algorithm. essentially, keep the best set of weights so far.




Classification boundary - PLA versus Pocket

PLA:



Pocket algorithm allows us to separate linearly separable data (it is basically separable but there are a few outliers which make it difficult to separate) - does a reasonable job at separating nonetheless

Pocket:



Outline

- Input representation
- Linear Classification
- Linear Regression regression \equiv real-valued output
- Nonlinear Transformation

Credit again

Classification: Credit approval (yes/no)

Regression: Credit line (dollar amount)

Input: $\mathbf{x} =$

age	23 years
annual salary	\$30,000
years in residence	1 year
years in job	1 year
current debt	\$15,000
...	...

Linear regression output: $h(\mathbf{x}) = \sum_{i=0}^d w_i x_i = \mathbf{w}^T \mathbf{x}$

The data set

Credit officers decide on credit lines:

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$$

$y_n \in \mathbb{R}$ is the credit line for customer \mathbf{x}_n .

Linear regression tries to replicate that.



How to measure the error

How well does $h(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$ approximate $f(\mathbf{x})$?

In linear regression, we use squared error $(h(\mathbf{x}) - f(\mathbf{x}))^2$

in-sample error:
$$E_{\text{in}}(h) = \frac{1}{N} \sum_{n=1}^N (h(\mathbf{x}_n) - y_n)^2$$

Illustration of linear regression



The expression for E_{in}

$$\begin{aligned} E_{\text{in}}(\mathbf{w}) &= \frac{1}{N} \sum_{n=1}^N (\mathbf{w}^{\text{T}} \mathbf{x}_n - y_n)^2 \\ &= \frac{1}{N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 \end{aligned}$$

where

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^{\text{T}} \\ \mathbf{x}_2^{\text{T}} \\ \vdots \\ \mathbf{x}_N^{\text{T}} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}$$

Minimizing E_{in}

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \|X\mathbf{w} - \mathbf{y}\|^2$$

$$\nabla E_{\text{in}}(\mathbf{w}) = \frac{2}{N} X^{\top} (X\mathbf{w} - \mathbf{y}) = \mathbf{0}$$

$$X^{\top} X \mathbf{w} = X^{\top} \mathbf{y}$$

$$\mathbf{w} = X^{\dagger} \mathbf{y} \quad \text{where} \quad X^{\dagger} = (X^{\top} X)^{-1} X^{\top}$$

X^{\dagger} is the 'pseudo-inverse' of X since $X^{\dagger} X = \text{Identity}$

The pseudo-inverse

d = number of parameters ($d+1$ due to bias or x_0)

N = number of training examples

$$\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$$

The diagram illustrates the dimensions of the matrices in the formula $\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$. It shows a large blue bracket on the left representing the matrix $(\mathbf{X}^\top \mathbf{X})^{-1}$, with a superscript -1 to its upper right. Inside this bracket, there is a smaller blue bracket representing the matrix $\mathbf{X}^\top \mathbf{X}$, with a dimension label $d+1 \times d+1$ below it. To the right of the large bracket is another blue bracket representing the matrix \mathbf{X}^\top , with a dimension label $d+1 \times N$ below it. A large horizontal blue bracket at the bottom spans the width of both the $(\mathbf{X}^\top \mathbf{X})^{-1}$ and \mathbf{X}^\top terms, with a dimension label $d+1 \times N$ centered below it.

The linear regression algorithm

- 1: Construct the matrix \mathbf{X} and the vector \mathbf{y} from the data set $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ as follows

$$\underbrace{\mathbf{X} = \begin{bmatrix} \text{---}\mathbf{x}_1^\top\text{---} \\ \text{---}\mathbf{x}_2^\top\text{---} \\ \vdots \\ \text{---}\mathbf{x}_N^\top\text{---} \end{bmatrix}}_{\text{input data matrix}}, \quad \underbrace{\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}}_{\text{target vector}}.$$

- 2: Compute the pseudo-inverse $\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$.
- 3: Return $\mathbf{w} = \mathbf{X}^\dagger \mathbf{y}$.

Chances in a real problem are overwhelming that \mathbf{X} will be an invertible matrix since the number of parameters/columns is usually small, with many rows (training examples).

Linear regression for classification

Linear regression learns a real-valued function $y = f(\mathbf{x}) \in \mathbb{R}$

Binary-valued functions are also real-valued! $\pm 1 \in \mathbb{R}$

Use linear regression to get \mathbf{w} where $\mathbf{w}^T \mathbf{x}_n \approx y_n = \pm 1$


In this case, $\text{sign}(\mathbf{w}^T \mathbf{x}_n)$ is likely to agree with $y_n = \pm 1$

Provide

Good initial weights for classification

which can then be fed to a classification algorithm (e.g. the 'pocket' algorithm) to get a better result. Not a good idea to only use linear regression since it will try to make all of the red points (see next slide) the value of -1, even the ones very far away from the theoretical perfect boundary between the groups - in doing so it is likely to move the line further downwards, perhaps meaning more blue points are incorrectly classified for the sake of already-red points (attempting to fit points irrelevant to classification, it can worsen the classification). Hence, better to use this as an initial weight for a proper classification algorithm.

Linear regression boundary




Outline


- Input representation
- Linear Classification
- Linear Regression
- Nonlinear Transformation

Linear is limited

Data:



Hypothesis:



Another example

Credit line is affected by 'years in residence'

but **not** in a linear way!

$[[...]]$ indicates that these return 1 if the condition is true and 0 if it is false

Nonlinear $[[x_i < 1]]$ and $[[x_i > 5]]$ are better.

Note the 'feature extraction/creation' seen here - a feature is made when you take the input and transform it into something else (a higher level representation of a raw input)

Can we do that with linear models?

Linear in what?

Linear regression implements

$$\sum_{i=0}^d \mathbf{w}_i x_i$$

Linear classification implements

$$\text{sign} \left(\sum_{i=0}^d \mathbf{w}_i x_i \right)$$

Algorithms work because of **linearity in the weights**

These algorithms are linear because they are linear in the weights - the inputs (x) are constants and the parameters/weights (w) are the variables that we care about

Transform the data nonlinearly

$$(x_1, x_2) \xrightarrow{\Phi} (x_1^2, x_2^2)$$

